

Forecasting Diabetes

Introduction to Problem & Data

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from IPython.display import HTML
```

```
# Load dataset for further analysis
```

```
df = pd.read_csv('https://raw.githubusercontent.com/Annieeee-Fang/Xianchen-s-Data-Bootcamp-Stuff/refs/heads/main/NHANE
```

```
# The url of the source
```

```
source_url = 'https://archive.ics.uci.edu/dataset/887/national+health+and+nutrition+health+survey+2013-2014+(nhanes)+
```

Preprocessing of the Dataset

```
# Extract the description of important variables from the webpage of the source
```

```
variables_table = pd.read_html(source_url)
```

```
df2 = variables_table[0]
```

```
df2
```

	Variable Name	Role	Type	Demographic	Description	Units	Missing Values
0	SEQN	ID	Continuous	NaN	Respondent Sequence Number	NaN	no
1	age_group	Target	Categorical	Age	Respondent's Age Group (senior/non-senior)	NaN	no
2	RIDAGEYR	Other	Continuous	Age	Respondent's Age	NaN	no
3	RIAGENDR	Feature	Continuous	Gender	Respondent's Gender	NaN	no
4	PAQ605	Feature	Continuous	NaN	If the respondent engages in moderate or vigor...	NaN	no
5	BMXBMI	Feature	Continuous	NaN	Respondent's Body Mass Index	NaN	no
6	LBXGLU	Feature	Continuous	NaN	Respondent's Blood Glucose after fasting	NaN	no
7	DIQ010	Feature	Continuous	NaN	If the Respondent is diabetic	NaN	no
8	LBXGLT	Feature	Continuous	NaN	Respondent's Oral	NaN	no
9	LBXIN	Feature	Continuous	NaN	Respondent's Blood Insulin Levels	NaN	no

```
# Create a dictionary of descriptions that are readable and short enough
```

```
new_descriptions = {
    "SEQN": "ID",
    "age_group": "Age_Group",
    "RIDAGEYR": "Age",
    "RIAGENDR": "Gender",
    "PAQ605": "Physical_Activity",
    "BMXBMI": "BMI",
    "LBXGLU": "Fasting_Glucose",
    "DIQ010": "Diabetes",
    "LBXGLT": "Oral",
    "LBXIN": "Insulin"
}
```

```
df = df.rename(columns=new_descriptions)
```

```
df
```



	ID	Age_Group	Age	Gender	Physical_Activity	BMI	Fasting_Glucose	Diabetes	Oral	Insulin
0	73564.0	Adult	61.0	2.0	2.0	35.7	110.0	2.0	150.0	14.91
1	73568.0	Adult	26.0	2.0	2.0	20.3	89.0	2.0	80.0	3.85
2	73576.0	Adult	16.0	1.0	2.0	23.2	89.0	2.0	68.0	6.14
3	73577.0	Adult	32.0	1.0	2.0	28.9	104.0	2.0	84.0	16.15
4	73580.0	Adult	38.0	2.0	1.0	35.9	103.0	2.0	81.0	10.92
...
2273	83711.0	Adult	38.0	2.0	2.0	33.5	100.0	2.0	73.0	6.53
2274	83712.0	Adult	61.0	1.0	2.0	30.0	93.0	2.0	208.0	13.02
2275	83713.0	Adult	34.0	1.0	2.0	23.7	103.0	2.0	124.0	21.41
2276	83718.0	Adult	60.0	2.0	2.0	27.4	90.0	2.0	108.0	4.99
2277	83727.0	Adult	26.0	1.0	2.0	24.5	108.0	2.0	108.0	3.76

2278 rows × 10 columns

```
df2['Description'].str.strip()
# Display the Original Variable Name, Descriptions that we are using, and the original Detailed Description
df2=df2.rename({'Variable Name': 'Original Variable Name', 'Description': 'Detailed Description'}, axis = 1)
df2['Description'] = df2['Original Variable Name'].map(new_descriptions)
pd.concat((df2['Description'], df2['Detailed Description']), axis = 1)
```



	Description	Detailed Description
0	ID	Respondent Sequence Number
1	Age_Group	Respondent's Age Group (senior/non-senior)
2	Age	Respondent's Age
3	Gender	Respondent's Gender
4	Physical_Activity	If the respondent engages in moderate or vigor...
5	BMI	Respondent's Body Mass Index
6	Fasting_Glucose	Respondent's Blood Glucose after fasting
7	Diabetes	If the Respondent is diabetic
8	Oral	Respondent's Oral
9	Insulin	Respondent's Blood Insulin Levels

```
#exchange 2's with 3's
df['Diabetes'] = df['Diabetes'].replace({2.0: 3.0, 3.0: 2.0})
df['Diabetes'].value_counts()
```



	count
Diabetes	
3.0	2199
2.0	58
1.0	21
dtype: int64	

```
df = df.drop('ID', axis=1)
```

```
df
```

	Age_Group	Age	Gender	Physical_Activity	BMI	Fasting_Glucose	Diabetes	Oral	Insulin
0	Adult	61.0	2.0	2.0	35.7	110.0	3.0	150.0	14.91
1	Adult	26.0	2.0	2.0	20.3	89.0	3.0	80.0	3.85
2	Adult	16.0	1.0	2.0	23.2	89.0	3.0	68.0	6.14
3	Adult	32.0	1.0	2.0	28.9	104.0	3.0	84.0	16.15
4	Adult	38.0	2.0	1.0	35.9	103.0	3.0	81.0	10.92
...
2273	Adult	38.0	2.0	2.0	33.5	100.0	3.0	73.0	6.53
2274	Adult	61.0	1.0	2.0	30.0	93.0	3.0	208.0	13.02
2275	Adult	34.0	1.0	2.0	23.7	103.0	3.0	124.0	21.41
2276	Adult	60.0	2.0	2.0	27.4	90.0	3.0	108.0	4.99
2277	Adult	26.0	1.0	2.0	24.5	108.0	3.0	108.0	3.76

2278 rows × 9 columns

✓ Preliminary Examination of the Dataset

✓ Descriptive Statistics

```
df['Diabetes'].value_counts()
```

	count
Diabetes	
3.0	2199
2.0	58
1.0	21

dtype: int64

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2278 entries, 0 to 2277
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age_Group             2278 non-null   object
1   Age                   2278 non-null   float64
2   Gender                2278 non-null   float64
3   Physical_Activity     2278 non-null   float64
4   BMI                   2278 non-null   float64
5   Fasting_Glucose       2278 non-null   float64
6   Diabetes              2278 non-null   float64
7   Oral                  2278 non-null   float64
8   Insulin               2278 non-null   float64
dtypes: float64(8), object(1)
memory usage: 160.3+ KB
```

```
df#.head(10)
```



	Age_Group	Age	Gender	Physical_Activity	BMI	Fasting_Glucose	Diabetes	Oral	Insulin
0	Adult	61.0	2.0	2.0	35.7	110.0	3.0	150.0	14.91
1	Adult	26.0	2.0	2.0	20.3	89.0	3.0	80.0	3.85
2	Adult	16.0	1.0	2.0	23.2	89.0	3.0	68.0	6.14
3	Adult	32.0	1.0	2.0	28.9	104.0	3.0	84.0	16.15
4	Adult	38.0	2.0	1.0	35.9	103.0	3.0	81.0	10.92
...
2273	Adult	38.0	2.0	2.0	33.5	100.0	3.0	73.0	6.53
2274	Adult	61.0	1.0	2.0	30.0	93.0	3.0	208.0	13.02
2275	Adult	34.0	1.0	2.0	23.7	103.0	3.0	124.0	21.41
2276	Adult	60.0	2.0	2.0	27.4	90.0	3.0	108.0	4.99
2277	Adult	26.0	1.0	2.0	24.5	108.0	3.0	108.0	3.76

2278 rows × 9 columns

```
df['Age_Group'].value_counts()
```



	count
Age_Group	
Adult	1914
Senior	364

dtype: int64

```
# The number of people of each ages tested
df['Age'].value_counts().sort_index(ascending=True)
```



	count
Age	
12.0	58
13.0	49
14.0	60
15.0	48
16.0	70
...	...
76.0	13
77.0	9
78.0	8
79.0	8
80.0	107

69 rows × 1 columns

dtype: int64

```
age_sorted_data = df.sort_values(by='Age', ascending=True)[df['Age_Group'] == 'Adult']
age_sorted_data
```

```
<ipython-input-15-c0978ea22d72>:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.  
age_sorted_data = df.sort_values(by='Age', ascending=True)[df['Age_Group'] == 'Adult']
```

	Age_Group	Age	Gender	Physical_Activity	BMI	Fasting_Glucose	Diabetes	Oral	Insulin
253	Adult	12.0	1.0	2.0	20.5	100.0	3.0	100.0	13.18
268	Adult	12.0	2.0	2.0	20.0	115.0	3.0	140.0	14.41
270	Adult	12.0	2.0	2.0	27.5	94.0	3.0	100.0	14.67
802	Adult	12.0	2.0	2.0	17.1	94.0	3.0	101.0	8.30
808	Adult	12.0	1.0	2.0	16.1	96.0	3.0	197.0	8.57
...
82	Adult	64.0	2.0	2.0	25.4	90.0	3.0	215.0	2.10
2214	Adult	64.0	2.0	2.0	23.7	110.0	2.0	82.0	9.65
403	Adult	64.0	1.0	2.0	27.5	99.0	3.0	119.0	4.47
948	Adult	64.0	2.0	2.0	41.5	91.0	3.0	149.0	15.52
1299	Adult	64.0	1.0	2.0	25.0	98.0	3.0	103.0	5.71

1914 rows x 9 columns

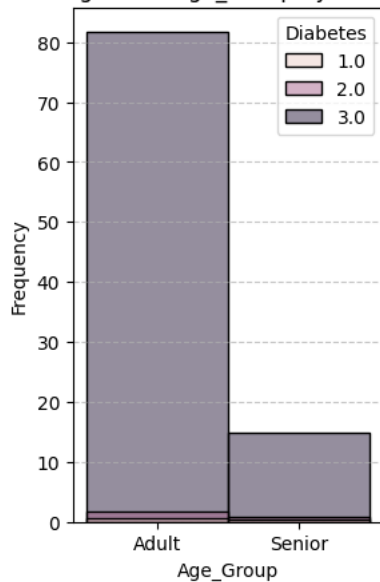
Initial Visualizations

```
variables1 = ['Age_Group', 'Gender', 'Physical_Activity', 'Oral']
```

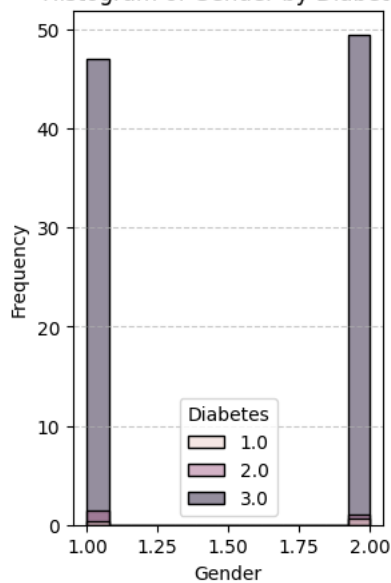
```
for col in variables1:  
    plt.figure(figsize=(3, 5))  
    sns.histplot(data=df, x=col, hue='Diabetes', kde=False, stat="percent")  
    t = f'Histogram of {col} by Diabetes'  
    plt.title(t)  
    plt.savefig(t+'.png')  
    plt.xlabel(col)  
    plt.ylabel('Frequency')  
    plt.grid(axis='y', linestyle='--', alpha=0.7)  
    plt.show()
```



Histogram of Age_Group by Diabetes



Histogram of Gender by Diabetes



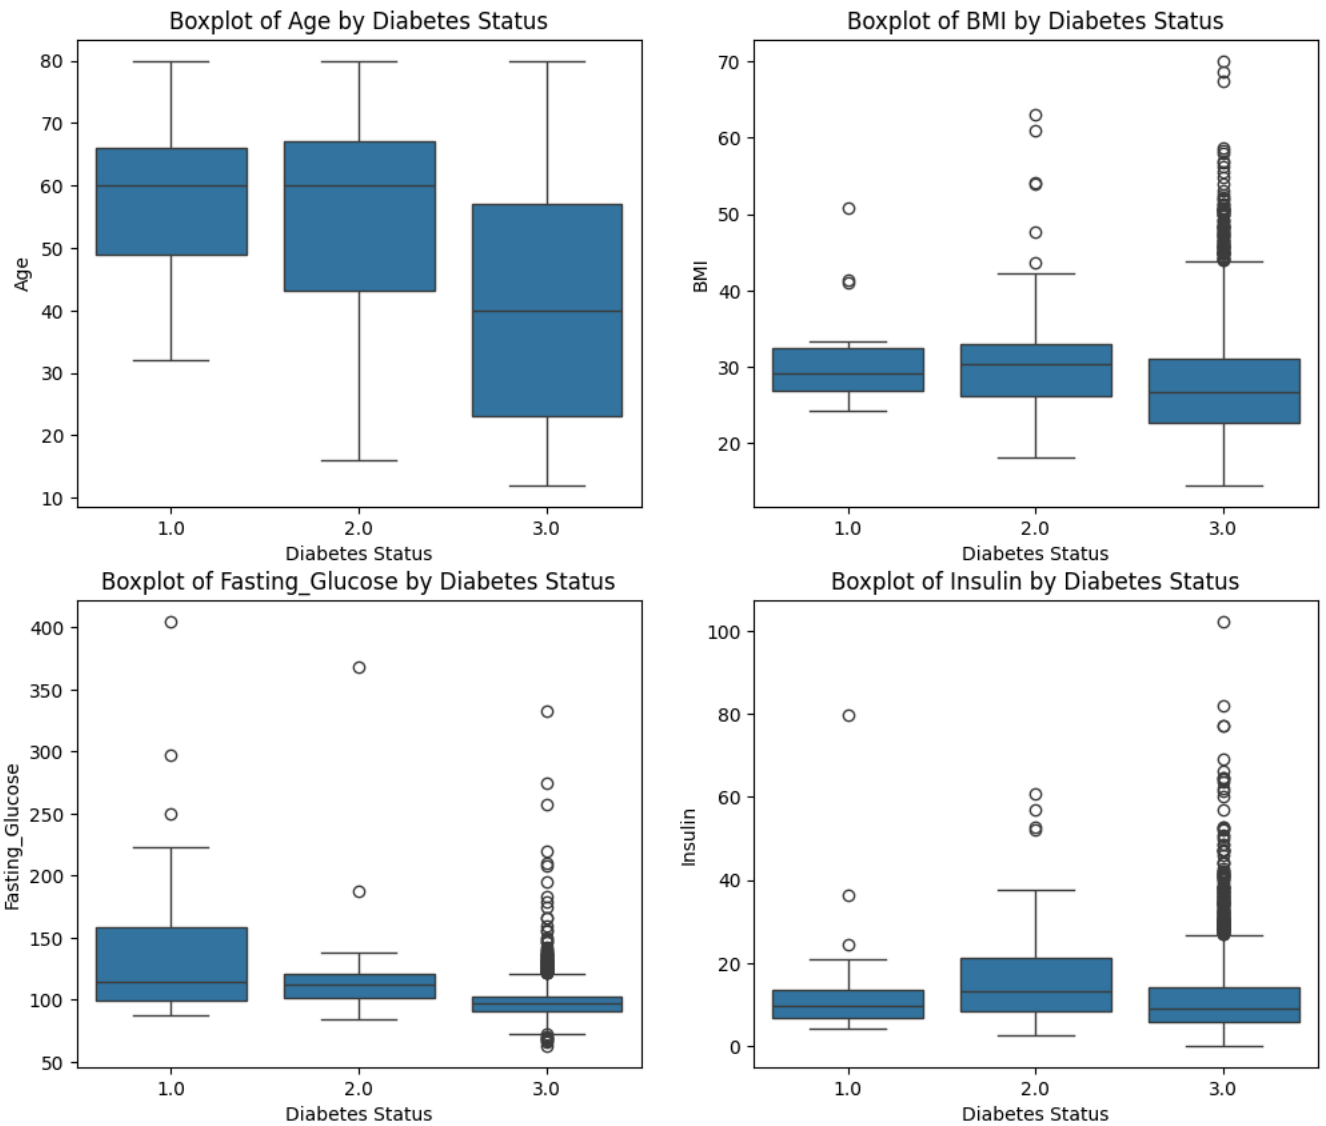
Histogram of Physical_Activity by Diabetes



```
variables2 = ['Age', 'BMI', 'Fasting_Glucose', 'Insulin']
fig, axes = plt.subplots(2, 2, figsize=(12, 10))
index=0
for i in range(2):
    for j in range(2):
        sns.boxplot(x='Diabetes', y=variables2[index], data=df, ax=axes[i][j])
        axes[i][j].set_title(f'Boxplot of {variables2[index]} by Diabetes Status')
        axes[i][j].set_xlabel('Diabetes Status')
        axes[i][j].set_ylabel(variables2[index])
        index+=1
plt.suptitle('Boxplots of Variables by Diabetes Status');
plt.savefig('Boxplots of Variables by Diabetes Status.png')
```



Boxplots of Variables by Diabetes Status



Modeling & Interpretations

```
X = df.drop(columns=['Diabetes'])
y = df['Diabetes']
```

Baseline Model

```
from sklearn.metrics import mean_squared_error
```

```
from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix
```

```
y.value_counts(normalize = True)
```



proportion	
Diabetes	
3.0	0.965320
2.0	0.025461
1.0	0.009219

dtype: float64

✓ Logistics Models

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import make_column_transformer
```

```
ohe = OneHotEncoder(drop='first')
sscaler = StandardScaler()
transformer = make_column_transformer((ohe, ['Age_Group']), remainder=sscaler)

lr = LogisticRegression(solver='lbfgs', max_iter=200)
pipe = Pipeline(['transform', transformer], ('logreg', lr))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

param_grid = [{'logreg__C': 10**np.linspace(-3, 3, 10)}]
```

```
from sklearn.metrics import make_scorer
```

```
from sklearn.metrics import precision_score #, accuracy_score, recall_score
```

```
methods_score = {}
new_precision = make_scorer(precision_score, average='macro', zero_division=1) # there existed some errors while us:
metrics= ["accuracy", new_precision]
for metric in metrics:
    lr_gridsearch_accuracy = GridSearchCV(pipe, param_grid, cv=10, scoring=metric)
    lr_gridsearch_accuracy.fit(X_train, y_train)
    score_accuracy = lr_gridsearch_accuracy.best_score_
    methods_score[metric] = score_accuracy
methods_score['precision'] = methods_score[new_precision]
```

```
methods_score.pop(new_precision)
print(methods_score)
```

```
➦ {'accuracy': 0.9648771993034287, 'precision': 0.9881092495846595}
```

✓ Decision Tree Regression Model

```
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score
from sklearn.inspection import DecisionBoundaryDisplay
```

```
ohe = OneHotEncoder(drop='first')
sscaler = StandardScaler()
transformer = make_column_transformer((ohe, ['Age_Group']), remainder=sscaler)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train_encoded = transformer.fit_transform(X_train)
X_test_encoded = transformer.transform(X_test)
```

```
train_scores = []
test_scores = []

for d in range(1, 20):
    dtree = DecisionTreeClassifier(max_depth=d, random_state=42)
    dtree.fit(X_train_encoded, y_train)


    y_train_preds = dtree.predict(X_train_encoded)
    y_test_preds = dtree.predict(X_test_encoded)

    train_scores.append(accuracy_score(y_train, y_train_preds))
    test_scores.append(accuracy_score(y_test, y_test_preds))
```

```
plt.figure(figsize=(10, 5))
plt.plot(range(1, 20), train_scores, '--o', label='Train Accuracy')
plt.plot(range(1, 20), test_scores, '--o', label='Test Accuracy')
plt.grid()
plt.legend()
plt.xticks(range(1, 20))
```

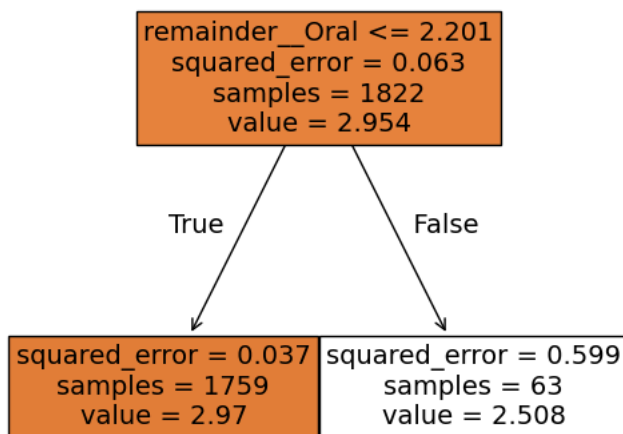


```
plt.xlabel('Max Tree Depth')
plt.ylabel('Accuracy')
plt.title('Decision Tree Depth vs. Test/Train Accuracy')
plt.savefig('Decision Tree Depth vs. Test and Train Accuracy.png')
```



```
from sklearn.tree import DecisionTreeRegressor
#fit a decision tree model with a max depth = 1 (lowest mse test score on graph)
dtree = DecisionTreeRegressor(max_depth = 1).fit(X_train_encoded, y_train)
```

```
plt.figure(figsize=(6,6))
plot_tree(dtree, filled=True, feature_names=transformer.get_feature_names_out().tolist(), fontsize=14)
plt.savefig('Decision Process.png');
```



```
max(test_scores)
```



```
0.9692982456140351
```

Random Forest Regression Model

```
from sklearn.ensemble import RandomForestClassifier
```

```
train_scores = []
test_scores = []

for n in range(10, 110, 10): # n is the numbers of estimators
```

```
rf = RandomForestClassifier(n_estimators=n, random_state=42)
rf.fit(X_train_encoded, y_train)
y_train_preds = rf.predict(X_train_encoded)
y_test_preds = rf.predict(X_test_encoded)
train_scores.append(accuracy_score(y_train, y_train_preds))
test_scores.append(accuracy_score(y_test, y_test_preds))
```